

FUNDING THE CAPABILITY-RESOURCE GAP: A MODEL OF OPTIMAL SCIENCE FUNDING

AYDIN MOHSENI, SIMON DEDEO, AND KEVIN J. S. ZOLLMAN

ABSTRACT. The grant funder’s challenge is to maximize the impact of a fixed budget under uncertainty about the researchers who might receive it. We develop a Bayesian model of this challenge: researchers differ in their epistemic capabilities and in resources; research output requires both, and funding compounds capability across rounds. We show that the optimal allocation funds the capability-resource gap: grants increase with capability and decrease with existing resources. Track records underdetermine the gap; peer review supplies an additional, noisy signal of capability. We find that this signal’s value increases with capability inequality and with its informativeness, and that under heavy-tailed capability distributions even a coarse signal suffices. We further find that uniform seed grants and lotteries, which forgo targeting, lose little output exactly where review is worth little, and that whom to fund matters far more than when to fund.

1. INTRODUCTION

Science funding is large and consequential. How best to allocate it remains poorly understood. The United States performed \$937 billion of R&D in 2023 and an estimated \$993 billion in 2024, roughly a fifth of it funded by the federal government (National Center for Science and Engineering Statistics 2026; National Science Board 2026). In spite of this, basic dimensions of allocation are contested. Whether funding should be concentrated on a few researchers or spread across many is a long-running dispute (Aagaard et al. 2020). Whether peer review is effective is disputed: among funded NIH grants, percentile scores predicted subsequent productivity barely better than chance (Fang et al. 2016), while other designs find modest but real predictive power (Gallo et al. 2014; Li and Agha 2015; Park et al. 2015). And whether review should be replaced, in part, by lotteries has moved from proposal to practice (Roumbanis 2019; Heyard et al. 2022). These disputes concern a single underlying problem: what a funder with a fixed budget and imperfect knowledge of researchers should do. This paper develops a

DEPARTMENT OF PHILOSOPHY, CARNEGIE MELLON UNIVERSITY
DEPARTMENT OF SOCIAL AND DECISION SCIENCES, CARNEGIE MELLON UNIVERSITY
DEPARTMENT OF PHILOSOPHY, CARNEGIE MELLON UNIVERSITY

E-mail addresses: amohseni@andrew.cmu.edu.

This work was supported by the John Templeton Foundation [grant number TODO].

model of that problem and works out its logic: whom to fund, whether to spend early or wait and learn, what information about researchers is worth, and what egalitarian alternatives lose.

Consider the challenge of the grant funder, whether a federal agency disbursing billions or a private foundation disbursing millions: you have a fixed budget and must choose how to allocate it across researchers and research projects to create the most positive impact. There are two intuitive schemes. Fund the track record: give to those who have produced the most, since past production is the best evidence of future production. Fund the under-resourced: give to those who have the least, since a dollar goes further where dollars are scarce. Each scheme can be intuitively compelling. Both, we will show, are suboptimal, and for the same reason: each accounts for only one of the two quantities the impact of a grant depends on.

We model this decision situation as follows. A funder with a fixed budget allocates grants, over multiple rounds, to researchers who differ in two ways: epistemic capability, what a researcher knows and can do, and resources, what a researcher has to work with. Research output requires both and is limited by the scarcer of the two; the funder can observe output, but not capability or resources themselves. A grant adds to a researcher's resources in the round it is given and, because doing research builds expertise, compounds their capability in later rounds. Each round, the funder observes the output that results, updates its beliefs about the researchers, and allocates again.

Our results are as follows. In the ideal condition where the funder has complete information regarding researchers' capabilities and resources, the optimal allocation funds the capability-resource gap: each funded researcher receives a grant increasing in their capability and decreasing in their existing resources, and the budget determines how many researchers receive funding at all. Each intuitive scheme fails because it accounts for only one of the two quantities. In reality the funder lacks direct access to the gap, and track records alone underdetermine it: we show that a funder relying on the track record performs far worse than one that estimates the gap.

Peer review supplies an additional signal of capability. Its value for a Bayesian funder is the expected value of its reduction of the funder's uncertainty regarding the gap, and that value is greatest where capability is unequally distributed and review informative; in fields with heavy-tailed distributions of capability, even a coarse signal captures most of this value. The same reasoning gives what the egalitarian alternatives lose: uniform seed grants and lotteries lose the value of

the targeting they displace, little where review is worth little, most where review is worth most. The budget is a variable throughout, so our results speak across funder scales: a small foundation, whose grants are modest beside researchers' total resources, is the funder for whom targeting matters most per dollar. Whether optimal funding should concentrate on a few researchers or spread across many depends on the size of the funder's budget relative to researchers' total resources. We show how strategic timing of when one provides funds matters far less than strategic targeting of whom one provides funds: how much of the budget to spend early rather than late shifts with the funder's budget relative to researchers' total resources, but the difference is small relative to the difference between funding the right and the wrong researchers.

Two substantive assumptions bound these results. Grants are consumed, a grant's durable residue (skills, publications, standing) being carried as capability; and capability and resources are complements throughout, neither fully substituting for the other. We defend each assumption where it is used.

We proceed as follows. §2 situates our model in the literatures on funding allocation, peer review, and lotteries. §3 presents our model. §4 derives the optimal allocation. §5 shows why track records and small grants cannot reveal the gap. §6 measures the value of peer review and relates it to the empirical evidence on review. §7 examines the timing of funding, including the strongest argument for spending early. §8 measures what seed grants and lotteries lose and asks how widely funds should be spread. §9 discusses implications and limitations, and §10 concludes.

2. RELATED LITERATURE

Our contribution sits at the intersection of four literatures: the empirical debate over concentrating or dispersing funds, the evidence on the efficacy of peer review, the case for funding lotteries, and formal models of science funding.

Whether funders should concentrate resources on a few researchers or spread them across many is a long-running dispute, and concentration is the current practice: the top 1% of NIH principal investigators receive close to 10% of research project grant funding, a share that has risen since 1985 (Katz and Matter 2020; Lauer and Roychowdhury 2021). Empirical studies of funded portfolios generally find diminishing marginal returns to funding per researcher, and reviews of this literature conclude that the evidence favors more dispersal than current practice (Mongeon et al. 2016; Aagaard et al. 2020). Concentration is also self-reinforcing: early grant winners accumulate advantages beyond the grant itself (Bol et al. 2018), a dynamic that formal models show can reward questionable research practices over

a career (Heesen 2024). This dispute typically concerns the returns to funding: whether output per researcher increases or decreases as grants concentrate. In our model, the answer also depends on the size of the funder’s budget relative to researchers’ total resources: with a budget large enough to relieve most bottlenecks, the optimal allocation itself reaches nearly every researcher, so little is lost by spreading funds widely; with a tight budget, effectiveness requires concentrating funds on the researchers with the largest capability-resource gaps (§8). The value of any targeting, in turn, depends on what the funder can know about researchers (§6).

The efficacy of peer review is contested. Among funded NIH grants, percentile scores predict subsequent productivity barely better than chance (Fang et al. 2016), and reviewers assessing the same application agree only weakly (Pier et al. 2018). Analyses with different designs find real but modest predictive power: better scores predict more publications, citations, and patents, including under controls for investigator history and institution (Li and Agha 2015), and a funding natural experiment finds higher-scored proposals outproduce lower-scored ones (Park et al. 2015). Both sides of this dispute measure review’s accuracy: how well scores predict outcomes among funded applicants. Our model instead measures review’s value: the expected value of its reduction of the funder’s uncertainty regarding the capability-resource gap. That value depends on the field; the same noisy signal can be worth a great deal in one field and nothing in another (§6), and no single number settles the question.

Doubts about review have made lotteries a live proposal. Advocates argue that randomizing among fundable applications saves review costs, reduces bias, and honestly reflects review’s limited ability to discriminate among fundable proposals (Fang and Casadevall 2016; Avin 2019a; Roumbanis 2019); several funders have implemented partial lotteries: the Health Research Council of New Zealand has allocated its Explorer Grants by lottery since 2013 (Liu et al. 2020), the Swiss National Science Foundation randomizes among ties (Heyard et al. 2022), and a German funding line runs a lottery before review (Luebber et al. 2025). The existing case measures what a lottery saves: Gross and Bergstrom (2019) model grant competition as a contest in which scientists invest costly effort in proposals, show that when only a small fraction of proposals is funded this effort can rival the scientific value of the research funded, and show that a partial lottery among proposals above a threshold removes the link between the wasted effort and the fraction funded. What has not been measured is what a lottery forgoes. Our model measures it: the value of targeting, which runs from negligible, where budgets are ample, signals uninformative, or capability evenly spread, to substantial, where

capability is heavy-tailed and budgets tight (§8). Together the two measures let a funder weigh the choice field by field.

Formal models of science funding have concentrated on the incentive costs of competition and on the value of exploration. Beyond the contest model above, Gross and Bergstrom (2025) find, in a game-theoretic model of contests for scarce rewards, that intensifying competition pushes scientists toward higher-risk, higher-return projects. Avin (2019b) simulates funding strategies on an epistemic landscape and finds that allocation with an explicitly random component significantly outperforms a peer-review-like strategy on large landscapes, because review concentrates effort on regions already known to be productive. Sorzano and Pueche-Granados (2026) treat allocation as a decision problem under heavy-tailed uncertainty about project value, pairing an empirical bibliometric signal with a biased-lottery mechanism. In each of these models, funding raises a single quantity, a project’s chance of success or a researcher’s productivity, and none treats the timing of spending as the funder’s choice. In our model, a researcher’s output increases in both capability and resources but is limited by the scarcer of the two, and the funder observes neither directly.

Our formulation and analysis of the capability-resource gap allows us to address questions raised by these literatures: how widely funds should be spread (§8), what peer review is worth (§6), and when lotteries are defensible (§8).

3. OUR MODEL

A funder faces a population of n researchers over T funding rounds. Researcher i has epistemic capability K_i and resources R_i . Initial capabilities and resources are drawn from Pareto distributions with tail parameters α_K and α_R ; smaller tail parameters produce more unequal fields, and a correlation parameter ρ allows capable researchers to tend to be well resourced. The funder knows these distributions; it never observes any researcher’s capability or resources directly.

In each round, researcher i produces research output¹ at random, in discrete units, with expected output²

$$\lambda_i = A \Lambda(K_i, R_i),$$

where A is a productivity constant and Λ is a production function that increases in both arguments but is limited by the scarcer of the two. Our main analysis

¹ We are agnostic about the interpretation of output units: they may be read as publications, as findings, or as units of scientific value.

² Output in each round is drawn from a Poisson distribution with rate λ_i , independently across researchers and rounds given the state.

takes Λ to be the harmonic mean:

$$\lambda_i = \frac{2AK_iR_i}{K_i + R_i}.$$

The harmonic mean makes capability and resources complements without making either an absolute requirement: a capable researcher with thin resources still produces, by working more slowly or borrowing equipment, but no surplus of one input fully substitutes for a shortage of the other. The harmonic mean sits between two boundary cases that fail in opposite directions: Cobb-Douglas, under which a capable researcher can compensate almost fully for vanishing resources, and Leontief, under which no compensation is possible at all. Appendix C repeats our analysis across the family between these boundaries, and §8 uses the family to ask how widely funds should be spread.

Capabilities increase through research. Between rounds,

$$K_i \leftarrow K_i + \epsilon \Lambda(K_i, R_i),$$

where ϵ is the compounding rate: doing research builds expertise, and the same bottleneck that limits output limits growth. Growth requires both inputs, and capability never declines. Resources, by contrast, do not accumulate: resources in a round are baseline plus grant, $R_i = R_{i0} + g_i$. The baseline R_{i0} recurs every round, representing the time and support a researcher has independent of the funder, while grants are consumed in the round they are given. This is a substantive assumption, not bookkeeping: whatever endures from a grant, the skills, publications, and standing it produces, is carried by the capability update, so funding's lasting effect flows through what researchers become rather than what they keep. Science is in this respect an accumulation process: success begets the means of further success.

The funder's total budget is fixed before any outcomes are observed. We parameterize it by a budget scale b : at $b = 0.5$, the purse equals the community's expected baseline resources for one round, and the purse scales linearly in b . (The exact normalization is in Appendix [X].) timing of spending, never its amount. Strategies that do not plan ahead spend an equal share of the purse each round. Forward-looking strategies may spend on any schedule: each round, they re-plan the entire remaining budget over the remaining horizon and execute the current round's allocation.

The funder learns as it goes. It observes each researcher's track record, the output they have produced to date, as it accumulates, and two further signals are available. A resource signal, observed once at the start, gives a noisy reading of each researcher's baseline resources (noise τ_R), as one might obtain from institutional

Strategy	Information	Planning
No funding	none	none
Uniform	none	none
Track record	track record	none
Myopic, without review	track record + resource signal	current round
Forward, without review	track record + resource signal	remaining horizon
Myopic, with review	track record + resource signal + review	current round
Forward, with review	track record + resource signal + review	remaining horizon

TABLE 1. The funding strategies compared.

information. A review signal, drawn afresh in each round that a strategy uses it, gives a noisy reading of capability (noise τ_K), as one might obtain from proposal evaluation and expert judgment. The review signal’s informativeness is governed by τ_K : the smaller the noise, the more informative the review. Bayesian strategies update beliefs about every researcher after each round and allocate on posterior expected marginal returns.

We compare seven funding strategies, which differ along two dimensions: information, with or without the review signal, and planning, myopic or forward-looking. A myopic strategy optimizes expected output in the current round; a forward-looking strategy optimizes expected output over the whole remaining horizon. Three non-Bayesian baselines complete the set: no funding, against which every result is normalized; uniform funding, which splits each round’s tranche equally; and track-record funding, which allocates in proportion to the previous round’s output. Uniform seed grants and lotteries, which modify these strategies rather than add to them, are introduced in §8.

Our conclusions rest either on analytic results supported by proofs or on simulation results that are established across the following parameter ranges.³ The tail parameters α_K and α_R range from 1.3 to 3.5. Distributions of research productivity are heavy-tailed: in Lotka’s classic data, the number of authors producing x papers falls off as x^{-2} , and later studies estimate exponents mostly between 2 and 3 (Lotka 1926; Nicholls 1989), corresponding to Pareto tail parameters near 1 to 2; our range covers this territory and extends to markedly more equal fields. The budget scale b ranges from 0.1 to 1, from a funder whose purse is small beside the community’s own resources to one able to relieve most researchers’ bottlenecks; §5 and §7 also examine deeper purses. The compounding rate ϵ ranges from near zero, where funding raises only current output, to 0.85, where this round’s funding substantially raises later capability. The review noise τ_K ranges from 0.05, nearly perfect review, to 20, nearly uninformative review, a range wide enough to contain

³ Unless a figure states otherwise, unstated parameters take the values $T = 2$, $n = 50$, $\epsilon = 0.1$, $b = 0.5$, $\alpha_K = \alpha_R = 2$, $\tau_K = \tau_R = 1$, and $\rho = 0$.

the noise implied by empirical estimates of review’s predictive accuracy (§6). The correlation ρ ranges from -0.5 to 0.8 , and the horizon T from one round to ten.

Randomness enters our model through the initial population draws, the signals, and realized output. Realized output shapes the funder’s beliefs and, through them, its allocations. Capability growth, by contrast, follows expected output: the update above uses Λ itself, so a researcher’s accumulation depends on their capability and resources, not on luck in any given round. This is a modeling choice: it treats a round as containing enough separate acts of research for luck to average out of accumulation, while leaving the funder’s inference fully exposed to noise. For the same reason, we score each strategy by the sum of expected outputs along its run, which equals realized output averaged over repetitions of the same allocations; this removes simulation noise from comparisons without favoring any strategy. Each configuration runs on 200 independently drawn populations, with common random numbers across strategies, so every comparison between strategies is paired. Throughout, we report each strategy’s improvement over no funding, in percent. The full specification, parameter table, and code are in Appendix [X].

With the model in place, the funder’s problem is exact: choose grants, given beliefs, to maximize expected total output over the horizon. Next we solve the case in which the funder has complete information regarding researchers’ capabilities and resources.

4. THE OPTIMAL ALLOCATION: FUND THE CAPABILITY-RESOURCE GAP

First, consider the simple case where the funder has complete information regarding researchers’ capabilities and resources. Further, we initially restrict our attention to a single round of funding, temporarily putting aside the contribution of resources to the accumulation of researchers’ capabilities across rounds. The funder’s problem is to divide the round’s tranche among researchers so as to maximize expected output:

$$\max_{g_1, \dots, g_n \geq 0} \sum_i A \Lambda(K_i, R_{i0} + g_i) \quad \text{subject to} \quad \sum_i g_i \leq \text{the round's tranche}.$$

Now, consider the marginal value of a dollar for such a funder. Under the harmonic production function,

$$\frac{\partial \lambda_i}{\partial g_i} = \frac{2A K_i^2}{(K_i + R_{i0} + g_i)^2}.$$

The marginal value of funding researcher i increases with their capability and decreases with the resources they already hold, including grants. An allocation is optimal exactly when no dollar can be moved to a researcher who can produce

more value with it, that is, when marginal values are equal across all funded researchers and no unfunded researcher's marginal value exceeds theirs.

Proposition 1 (the gap rule). *For every researcher, the optimal grant fills the gap, if any, between the resources they hold and a target proportional to their capability:*

$$g_i^* = \max(c K_i - R_{i0}, 0),$$

where the constant $c > 0$ is set by spending the budget.⁴

We can understand the rule as follows. Each researcher's target is c times their capability, and their gap is the difference between their target and the resources they hold. The rule grants every researcher with a positive gap exactly their gap, and grants nothing to the rest; the constant c is defined as the value at which these grants sum to the budget. Because c is defined by the budget, the budget always covers the gaps: a question of which gaps to fund first does not arise. When the budget is small, c is small: targets are low, few researchers have positive gaps, and their grants are small. When the budget is large, c is large: targets are high, and funding extends across nearly the whole population. In both cases the funded researchers, those with positive gaps ($R_{i0} < c K_i$), are those with the least resources relative to their capability, and these are the researchers for whom a dollar has the greatest marginal value.⁵ The condition $R_{i0} < c K_i$ draws a frontier through the population, the funding frontier, separating the funded from the unfunded (Figure 1).

The gap rule explains why each of the intuitive funding schemes discussed earlier fails. Funding the track record fails because output reflects resources as well as capability: a productive researcher whose resources approach or exceed their target has a small or negative gap, and an additional dollar adds little output there. The failure is not merely a shortfall from the optimum: because track-record funding channels money toward researchers whose bottleneck is not resources, it can produce less output than spreading the budget evenly, that is, than not targeting at all.⁶ Funding the under-resourced fails for the mirror-image reason: scarcity is no evidence of capability, and dollars given to researchers of low capability add little output however scarce their resources. It, too, can produce less than not targeting

⁴ $c = \sqrt{2A/\nu} - 1$, with ν the common marginal value at the optimum. The proof is in Appendix A.

⁵ The optimality condition behind the rule, funding to a common marginal value, is the water-filling solution familiar from information theory (Cover and Thomas 2006). If all researchers shared one capability, the rule would reduce to filling each researcher's gap below one common target, the poverty-gap transfer of development economics (Foster et al. 1984); our rule scales the target with capability.

⁶ Example 1 in Appendix A makes this exact with two researchers.




FIGURE 1. The gap rule in one population. Each point is a researcher, positioned by their resources and capability. The line is the funding frontier, $R = cK$, here at budget scale $b = 0.1$. The rule funds every researcher whose resources fall below the frontier (filled points), and each grant, drawn as an arrow, moves its researcher exactly to the frontier: the arrow's length is the researcher's gap, $cK_i - R_{i0}$. Researchers at or beyond the frontier (open points) receive nothing.

at all.⁷ Each scheme accounts for one of capability and resources, whereas the marginal value of a dollar is determined by how a researcher's resources compare with their capability, and so responds to both.

How much targeting matters depends on the budget. A researcher's expected output increases with every dollar of resources, but with diminishing marginal returns, and it never exceeds $2AK_i$: expected output is bounded above by a quantity set by capability, which it approaches but does not attain. When the budget is large, optimal and uniform funding therefore differ little in the output they produce: under either, every researcher's output is near its upper bound, and the difference between the two vanishes as the budget increases (Corollary 2, Appendix A). When the budget is small, dollars are scarce, researchers differ in the marginal value of a dollar, and output depends on where each dollar goes. Targeting matters where money is tight and little where it is ample (Figure 2). This comparative static recurs: it sets how much output is lost by overriding the optimal allocation with seed grants and lotteries (§8), and it determines when optimal funding concentrates rather than spreads (§8).

⁷ Again, Example 2 in Appendix A makes this exact.




FIGURE 2. The value of targeting decreases in the budget. The value of targeting is the additional expected output from the optimal allocation over uniform funding, as a fraction of uniform funding's gain over no funding; the budget scale b is the budget as a fraction of the population's total baseline resources. The curve shows one population of $n = 75$ researchers drawn from the model's default distributions ($\alpha_K = \alpha_R = 2$, $\rho = 0$, $A = 1$); the value of targeting decreases toward zero as the budget increases (Corollary 2).

To use the gap rule, however, a funder requires an estimate of the gap between researchers' capabilities and their resources. Next we turn to how such estimation can work: first from the track record alone, then with peer review added.

5. WHAT THE TRACK RECORD CAN AND CANNOT SUPPLY

The funder observes track records: each researcher's realized output, round by round. A track record is evidence about the capability-resource gap, and the question for this section is how far that evidence goes.

Consider first what output can identify in principle. Expected output is one number, produced by capability and resources jointly: every pair of capability and resources with $2AKR/(K + R) = \lambda$ produces the same expected output λ . A researcher of high capability and modest resources and a researcher of modest capability and abundant resources can therefore produce identical records. Output does carry some information about capability, a lower bound: since expected output never exceeds $2AK$, a researcher who produces λ has capability above $\lambda/2A$. Beyond that bound, the record cannot separate the two researchers above, and these are exactly the researchers the gap rule treats most differently: the first may have the largest gap in the population, the second may have none. Cross-sectional output, however carefully recorded, underdetermines what the funder most needs to know.

Our model measures how much output this underdetermination loses. A Bayesian funder relying on records alone allocates far from the gap rule, and further rounds of records improve its allocation only slowly; its output falls well short of the complete-information benchmark.⁸ This shortfall is the span from records alone to complete information. Next we ask how much of that span a review signal recovers, and there display this records-only baseline beside a funder holding the review signal.

The record is least informative exactly where the gap is largest. When a researcher’s resources are small relative to their capability, expected output is approximately $2AR$: proportional to resources, nearly independent of capability (Figure 3). On a thin grant, a researcher of capability 1 and a researcher of capability 100 produce nearly indistinguishable output. Small grants therefore yield almost no information about who is capable; separating capability from resources by observing output requires funding depth, grants comparable to capability itself. In our simulations of a resource-poor community with no review signal, a funder allocating at standard depth gains almost nothing over uniform funding, while the same funder with deep grants recovers most of the value of discrimination (§7 measures this and its consequences for timing). For nascent and resource-poor fields the implication is direct: thin grants cannot reveal who is capable.

What the record cannot supply, then, is a signal of capability separate from resources. That is what peer review can provide under the right conditions. Next we measure what such a signal is worth, and state the conditions.

6. PEER REVIEW: WHAT A CAPABILITY SIGNAL IS WORTH

Peer review enters our model as a second source of evidence about researchers: a noisy signal of each researcher’s capability. The value of review, for a Bayesian funder, is the expected value of its reduction of the funder’s uncertainty regarding the capability-resource gap. We can quantify this value in terms of output: the additional expected output of a funder allocating with the review signal over the same funder allocating without it.

The direction of review’s effect is what one would expect: a more informative signal brings the funder’s allocation closer to the optimal allocation and its output closer to what complete information would deliver, and with a sharp enough signal the funder recovers most of the output that a funder relying on records alone leaves

⁸ At the default parameters, the correlation between its grants and the optimal grants is 0.08 in the first round and reaches only 0.41 by round twenty; its output falls short of the complete-information benchmark by 11.4 percent of no-funding output (24.6 percent under heavy-tailed capability).




FIGURE 3. Thin grants are uninformative about capability. Expected output as a function of resources, for researchers of capability $K = 3, 10$, and 30 ($A = 1$). Every curve has slope $2A$ at the origin: where resources are small relative to capability, the curves coincide, and researchers of very different capabilities produce nearly identical output. The curves separate where resources are comparable to capability: observing output can identify capability only at depth.

unrealized.⁹ Two further effects are less obvious: (1) a funder relying on records alone does not catch up, and (2) what a review signal is worth varies widely with the field. The first effect has a structural cause. Capabilities compound across rounds while resources do not accumulate, so each researcher’s output approaches the level their resources sustain, and output at that level carries less and less information about capability. Further rounds of records do improve the funder’s allocation, but slowly: after twenty rounds of records the funder’s allocation remains further from the optimal allocation than a funder with the review signal is after one (Figure 4).

Our main result delineates how much review is worth as a function of the field (Figure 5). Review is worth most where capability is heavy-tailed and the budget is tight, and in the heavy-tailed case even a signal of moderate informativeness captures most of its value. The reason is twofold. Where capability is heavy-tailed, most of the output a funder can add comes from getting resources to

⁹ At the default parameters, the correlation between the funder’s grants and the optimal grants increases from 0.22 at the noisiest signal in our range to 0.94 at the sharpest; under heavy-tailed capability ($\alpha_K = 1.3$), from 0.34 to 0.96. The output shortfall to the complete-information benchmark decreases from 10.4 to 1.6 percent of no-funding output; under heavy-tailed capability, from 17.2 to 1.7 percent. The increase in review’s value is not strictly monotone at the sharpest signals: under heavy-tailed capability, its value at $\tau_K = 0.05$ is about two percent below its value at $\tau_K = 0.3$, a small but statistically solid reversal (paired across seeds, $z = 4.5$). The cause is the funder’s other estimate: resources are also observed with noise, and a slightly noisy capability signal tempers the allocation against errors in the resource estimate, while a near-perfect one commits the allocation to them. With the resource signal nearly noiseless, the sharper capability signal is again worth more (paired $z = 5.8$).




FIGURE 4. Records do not catch up. The correlation of each funder’s grants with the optimal grants, round by round, at the default parameters ($\tau_K = 1$, fifty seeds, twenty rounds). The records-only funder improves slowly, from 0.08 in the first round to 0.41 by round twenty, and remains further from the optimal allocation than the funder with the review signal is in round one (0.81). Under heavy-tailed capability the pattern is the same: 0.07 to 0.43, against 0.92.

the few researchers of unusually high capability. And precisely those researchers are the easiest to identify: they stand far apart from everyone else, so even a noisy signal separates them. Identifying the most capable researchers captures most of the value and requires the least information. Sharpening review beyond a modest level therefore adds little.¹⁰ This ordering of review’s value by field does not depend on our production form: it holds across the family of technologies from Cobb-Douglas to Leontief (Appendix C).

A corollary concerns where review effort goes. In our model, what adds output is separating the exceptional researchers from the rest, and that separation requires little information; fine distinctions among the many applications of comparable middling merit add comparatively little output. A review system that invests most of its effort in exactly those fine distinctions may be misallocating it.

¹⁰ At three times the default signal noise, review retains 82 percent of its maximum value under heavy-tailed capability ($\alpha_K = 1.3$), 44 percent at the default distribution ($\alpha_K = 2$), and 10 percent where capability is spread more evenly ($\alpha_K = 3.5$). The noise level at which review’s value falls to half its maximum is roughly $\tau_K = 10$, 2.5, and 0.75 in the three cases; the default noise is $\tau_K = 1$, and our sweep runs from near-perfect review ($\tau_K = 0.05$) to nearly uninformative review ($\tau_K = 20$). Fifty seeds per cell.




FIGURE 5. The value of review by field. Each curve shows the share of the records-only shortfall that review recovers: the additional expected output of a funder allocating with the review signal over the same funder allocating without it, as a fraction of the output the records-only funder leaves unrealized relative to complete information. Signal noise τ_K increases to the right, so informativeness decreases; the default is $\tau_K = 1$. Under heavy-tailed capability ($\alpha_K = 1.3$), review recovers most of the shortfall across a wide range of noise; where capability is spread more evenly ($\alpha_K = 3.5$), review recovers little at any noise. Fifty seeds per point; the records-only shortfall is 24.6, 11.4, and 3.5 percent of no-funding output at $\alpha_K = 1.3$, 2, and 3.5.

The remaining comparisons run the other way. Where capability is spread more evenly, no researcher matters much more than another, so review of any informativeness adds little.¹¹ And where the budget is ample, optimal and uniform funding produce nearly the same output, so information about whom to fund has little left to add: review is worth most where money is tight.

These results bear on the dispute over the efficacy of peer review. Among funded NIH grants, percentile scores predict subsequent productivity barely better than chance. In our model, a signal of that accuracy is very noisy review.¹² Whether such review is worth having depends on the field. At the noisiest signal in our sweep, review still recovers about a third of the records-only shortfall under heavy-tailed capability, under a tenth at the default distribution, and essentially none

¹¹ At $\alpha_K = 3.5$, the entire records-only shortfall is 3.5 percent of no-funding output (24.6 percent under heavy-tailed capability), and even near-perfect review recovers only 72 percent of it: review's value at its maximum is 2.5 percent of no-funding output, against 23 percent under heavy-tailed capability.

¹² An AUC of 0.54 (Fang et al. 2016) corresponds to $\tau_K \gtrsim 20$ at the default parameters, the noisiest signal in our sweep. The mapping from measured accuracy to signal noise itself depends on the field: at the same noise, accuracy is higher under heavy-tailed capability, because the most capable researchers stand further apart.




FIGURE 6. Overtrusting review loses more output than undertrusting it. The value of review as a percent of no-funding output, as the funder’s believed signal noise varies while the true noise is $\tau_K = 3$ (dotted rule); bars show one standard error over two hundred seeds. Value peaks where belief matches truth. To the right of the peak, undertrust forgoes part of the value; to the left, overtrust treats a noisy signal as sharp, and at the default capability distribution ($\alpha_K = 2$) it takes the value below zero: worse than ignoring review. Under heavy-tailed capability ($\alpha_K = 1.3$) overtrust reduces the value without reversing it.

where capability is spread more evenly.¹³ Accuracy is one input to review’s value; the field’s capability distribution and the funder’s budget are the others. No single number settles the dispute: review of the same informativeness is worth several times more in one field than in another, and the same measured accuracy corresponds to different informativeness in different fields.

One condition qualifies all of this: the funder must weight review at its true informativeness. A funder that overtrusts review, treating a noisy signal as sharp, can do worse than one that ignores review altogether: with moderately spread capability, overtrust turns review’s value negative; with heavy-tailed capability it reduces the value without reversing it (Figure 6).¹⁴ Overtrusting review loses more output than undertrusting it.

¹³ 31, 9, and 0.1 percent of the records-only shortfall at $\tau_K = 20$, for $\alpha_K = 1.3$, 2, and 3.5; fifty seeds per cell.

¹⁴ Each negative cell in Figure 6 is within two standard errors of zero, but all are negative. The tenfold comparison, across true noise levels: a funder that overtrusts tenfold (true $\tau_K = 3$, belief 0.3) loses more than review’s full calibrated value at the default capability distribution, while a funder that undertrusts tenfold (true $\tau_K = 0.3$, belief 3) retains over half of it.

Review can inform whom to fund. Next we turn to when to fund: how funding should be spread across rounds, and how much that choice matters.

7. THE TIMING OF FUNDING

Funding in our model is spread over rounds. The funder holds a fixed budget for the whole horizon, and between rounds two things move: researchers produce output, which the funder observes, and capabilities compound through the research that output represents. Grants themselves are consumed within their round. Up to now the funder has spent its budget in equal parts each round; this section asks whether it should not: whether to spend early, evenly, or late, and how much the choice matters.

The strongest argument for spending early concerns resource-poor fields. Researchers without resources produce nothing; where nothing is produced, nothing is observed and nothing compounds; so money held for later rounds produces neither evidence nor growth in the meantime. On this argument a funder facing a poor field should spend at once, to start the engine. The argument fails on inspection. A funder that spends everything in the first round allocates before observing any output, exactly as blind as one that spends everything in the last; in the extreme case of a field with no resources at all, the two schedules produce identical expected output. What the argument does establish is a complementarity: early spending enables observation and growth, and late spending exploits what early spending revealed. That favors putting some money early. It does not favor putting most of it early.¹⁵

Our simulations bear this out and say which way the asymmetry runs: observation comes first, money after. Wherever capabilities compound at more than a token rate, the optimal schedule holds a below-even share in the first round and deploys the bulk once output has revealed who is capable; spending late is rewarded because late money is informed money, and because early output compounds on its own (Figure 7). Resource poverty raises the early share and decreases the intensity of the effect, but does not reverse it.¹⁶

How much a funder can learn from its own grants depends on their size. A researcher on a thin grant produces output nearly proportional to the grant and

¹⁵ In a planner-free check over a grid of two-block schedules executed by a funder re-deciding each round, the best schedule that puts an above-even share in the early rounds loses to the even schedule, by 1.6 percent of the even schedule's output at negligible compounding and by 3.5 percent at the default compounding rate (six rounds, resource-poor community, two hundred seeds).

¹⁶ Figure 7 shows three of the five baseline resource scales in the map; the other two lie between the curves shown. Every cell whose schedule leans early has a compounding rate $\epsilon \leq 0.03$, and there the schedule's center of mass is within 0.005 of even.




FIGURE 7. Compounding, not poverty, sets the schedule’s tilt. The schedule’s center of mass (0.5 is even; higher is later) as the compounding rate ϵ increases, for communities of baseline resource scale $r_{\min} = 3, 0.3$, and 0.001 . Below $\epsilon \approx 0.03$ every schedule is within half a percent of even; as compounding increases, every community’s schedule leans later. Resource poverty raises the early share and decreases the intensity of the effect, but does not reverse it. Two hundred seeds per cell.

nearly independent of capability, so a funder whose grants are thin learns almost nothing it can exploit later: in a resource-poor field with no review signal, such a funder gains almost nothing over uniform funding. Deep grants change this. As grants approach the scale of capability itself, output separates the capable from the rest, and the funder recovers most of the value of discrimination, on a schedule that puts observation before money: a small early share to make output observable, the mass deployed once informed (Figure 8).¹⁷

The pattern has a single main cause. Capabilities compound through all research output, funded or not; grants add to output and so add to compounding. Separating the two channels: with only the grant-fed channel active, the optimal schedule leans mildly early, since earlier grants compound longer. With only the free channel active, it leans strongly late. The free channel dominates once it reaches even a fraction of the grant-fed rate, so in any field where research compounds regardless of who is funded, the funder’s reason to wait outweighs its reason to hurry.¹⁸

¹⁷ In a resource-poor community with heavy-tailed capability, no review signal, and six rounds (two hundred seeds): the informed funder’s gain over uniform funding increases from a third of one percent of uniform funding’s output at standard depth to nine percent at six times that depth and ten percent at twelve times, roughly the gain a sharp review signal delivers at standard depth (nine to eleven percent). In these runs the budget is scaled against capability rather than resources: at the six-times depth, the average grant per researcher per round is half of mean capability.

¹⁸ With the free channel off, the schedule’s center of mass is 0.48 to 0.49 across grant-fed rates, strictly below even; with the grant-fed channel off, 0.53 to 0.68. The free channel’s rate takes over at roughly one eighth to one third of the grant-fed rate (two hundred seeds, five rounds).




FIGURE 8. Observation first, money after. The share of the budget the funder spends in each of six rounds, in a resource-poor field with heavy-tailed capability and no review signal. At standard depth (filled marks) the funder learns nothing from its own thin grants and the optimal schedule is even. With deep grants (open marks) the funder holds the first-round share below even and deploys the bulk once output has revealed who is capable.

Finally, what is the schedule choice worth? Less than information, in every setting we test. Across all horizons and compounding rates in our sweeps, the gain from deliberate scheduling over even installments re-decided each round stays below the review signal's value: at the default compounding rate and below it is smaller by an order of magnitude or more, and it approaches the signal's value only at the joint extreme of the highest compounding rate and the longest horizon we consider (Figure 9).¹⁹ The gain increases roughly with the square of the horizon over short horizons and saturates beyond them.²⁰ And it exists only where capability and resources are complements, as in our production form: under Cobb-Douglas production the value of scheduling is indistinguishable from zero and the optimal schedule stays within about a hundredth of even (Appendix C, Table 2). For a funder thinking about how to maximize their impact, choosing whom to fund outweighs choosing when.

Next we turn to two alternatives to targeting that recur in funding policy: uniform seed grants and lotteries.

¹⁹ As a fraction of no-funding output, the scheduling gain ranges from indistinguishable from zero to 0.7 percent across the thirty-two tested cells; the review signal's value in the same cells ranges from 0.8 to 7.5 percent.

²⁰ Fitted exponents 2.3 and 2.1 for horizons up to five rounds at $\epsilon = 0.3$ and 0.85; 1.4 and 0.8 over five to ten rounds.




FIGURE 9. Choosing whom outweighs choosing when. The gain from deliberate scheduling (a funder that plans its budget’s spread over the horizon, over one that spends even installments and re-decides each round) as a share of the review signal’s value, by horizon and compounding rate ϵ ; the default rate is 0.3. In every setting we test the gain is below the signal’s value: an order of magnitude below or more at the default rate and below, and approaching the signal’s value only at the joint extreme of the highest rate and the longest horizon. The two lowest curves, at $\epsilon = 0.15$ and 0.05, and the curve at 0.55 were swept to five rounds.

8. SPREADING FUNDS: SEED GRANTS, LOTTERIES, AND CONCENTRATION

Targeting sends each dollar to the researcher expected to produce the most with it. Two alternatives to targeting recur in funding policy, and each has a rationale. Uniform seed grants give every researcher the same amount, so that a researcher whom review misjudges still receives something; in practice a funder gives out part of its budget as seed grants and targets the rest. Lotteries give grants to applicants chosen at random, usually from among those who pass a review threshold; their advocates argue that review cannot reliably rank the applicants it deems fundable, that review is costly, and that review is biased. Both alternatives give up part of targeting on purpose. Our model can say how much expected output is lost by giving up targeting, field by field, and where the loss is small. What the alternatives save, in the cost of review and in the effort researchers spend on proposals, lies outside our model and has been studied by others;²¹ we return to the comparison at the close of the section.

²¹ Gross and Bergstrom (2019) model grant competition as a contest in which scientists spend costly effort on proposals. They show that when only a small fraction of proposals is funded, this effort can rival the scientific value of the research funded, and that a lottery among proposals above a quality threshold removes the link between the wasted effort and the fraction funded.

Both alternatives lose the same thing: the value of targeting, the additional expected output that the optimal allocation produces over an equal division of the same budget. How much output is lost by giving up targeting depends on the field. For a funder allocating with records and review, targeting adds more than uniform funding’s entire gain over no funding in a heavy-tailed field with sharp review and a tight budget, and about a quarter of uniform funding’s gain in the default field at the largest budget we sweep.²² Giving up targeting loses little output under the same conditions that make review worth little: evenly spread capabilities, large budgets, uninformative peer review.

Consider uniform seed grants first. The output lost increases faster than the fraction of the budget given out as seed grants: giving out half of the budget as seed grants loses more than twice the output that giving out a quarter loses (Figure 10). The size of the loss is set by what targeting is worth in the field. When a review-informed funder gives out three quarters of its budget as seed grants, expected output falls by about six percent of no-funding output in a heavy-tailed field with sharp review, and by about one percent in the default field.²³ When the whole budget is given out as seed grants, the scheme is uniform funding, and the full value of targeting is lost. Whether the reasons for seed grants justify the loss is a judgment our model does not make; the model supplies the size of the loss.

Now consider lotteries. A lottery gives each applicant in its pool a chance at a full grant instead of a certain smaller one, and on average the applicant receives the same amount either way. A researcher’s expected output increases with resources at a diminishing rate, so a certain grant produces more expected output than a chance at a larger grant of the same average size. Dividing the pool’s money equally among the pool therefore produces more expected output than any lottery over the same pool, and a lottery over all applicants produces less expected output than uniform funding (Proposition 2, Appendix B). To measure how much less, and how lotteries compare with selection by review, we evaluate five schemes in a single round (Figure 11). The first scheme divides the budget equally among the fifth of researchers with the highest review scores. The second divides the budget equally among the top two fifths. The third is the partial lottery that some funders run: it funds half of the top two fifths at random, with the same

²² Over two rounds: the review-informed funder’s gain over uniform funding, as a fraction of uniform funding’s gain over no funding, decreases from 1.17 to 0.72 across budget scales $b = 0.1$ to 1 with heavy-tailed capability and sharp review ($\tau_K = 0.3$), and from 0.74 to 0.25 in the default field with the default signal. The two regimes differ in both the capability distribution and the signal’s noise. Two hundred seeds per cell.

²³ Across the swept cells, the output lost is approximately the informed funder’s gain over uniform funding, times the fraction of the budget given out as seed grants, times a factor between 0.10 and 0.86; the factor is above 0.3 at tight budgets and in heavy-tailed fields. Two hundred seeds per cell.




FIGURE 10. The output lost to uniform seed grants. A review-informed funder gives out a fraction of its budget as equal grants to all researchers and targets the rest. The figure shows the expected output lost, as a percent of no-funding output, over two rounds. The loss increases faster than the fraction given out as seed grants. The size of the loss is set by what targeting is worth in the field: the loss is largest in a heavy-tailed field with sharp review (upper curves, $\tau_K = 0.3$) and small where targeting is worth little (lowest curve: the default field and the default review signal). Two hundred seeds per point.

grant size as the first scheme. The fourth is uniform funding. The fifth is a lottery over all researchers, again with the first scheme's grant size. In a heavy-tailed field with a tight budget, selection by review, the first scheme, gains about three quarters of what the optimal allocation gains when review is sharp and about half when review is very noisy. The partial lottery gains less than selection by review, and most of the difference comes from ignoring the review ranking within the top two fifths, not from the randomness itself.²⁴ The lottery over all researchers gains least, less than uniform funding. In an evenly spread field with a large budget, the ordering reverses: uniform funding gains more than four fifths of what the optimal allocation gains, and every scheme that concentrates grants, whether by review score or by chance, gains less.

When does a lottery among the fundable lose little output? When the review ranking it ignores is worth little. Replacing the review ranking within the top two fifths by chance loses a few hundredths of the optimal allocation's gain in

²⁴ At review noise $\tau = 1$: selection by review gains 0.75 of what the optimal allocation gains, the equal division among the top two fifths 0.61, and the partial lottery 0.56. Of the 0.18 by which the partial lottery falls short of selection by review, ignoring the ranking accounts for 0.14 and the randomness for 0.04. Single round, a fifth of the population funded, two thousand populations per point.




FIGURE 11. What five funding schemes gain, by field. Each scheme's gain over no funding as a share of the optimal allocation's gain, in a single round with a fifth of the population funded, as review noise increases. Filled circles: an equal division among the top fifth of researchers by review score. Open circles: an equal division among the top two fifths. Triangles: the partial lottery, which funds half of the top two fifths at random. Dashed rule: uniform funding. Dotted rule: a lottery over all researchers. In a heavy-tailed field with a tight budget (left), selection by review gains most, and its gain declines slowly as review grows noisier; the partial lottery gains less, mainly because it ignores the review ranking within the top two fifths; the lottery over all researchers gains least, less than uniform funding. In an evenly spread field with a large budget (right), the ordering reverses: uniform funding gains most, and every scheme that concentrates grants gains less. Two thousand populations per point.

an evenly spread field or when review is very noisy, and about a sixth of that gain in a heavy-tailed field with sharp review.²⁵ These are the fields in which fine distinctions among applicants are worth little. The lottery debate rarely asks a further question, whether to concentrate grants at all. In an evenly spread field, our model's answer is to spread the budget thinly across many researchers, not to hold a lottery among a fundable few. Whether a lottery's savings in review costs and proposal effort outweigh the output it loses depends on both quantities. The contest models cited above measure the savings; our model measures the loss.

Seed grants and lotteries both spread funds more widely than targeting does, and behind both lies the older question of whether funding should be concentrated on

²⁵ The partial lottery's shortfall relative to selection by review: 0.01 to 0.07 of the optimal allocation's gain in the evenly spread field, and in the heavy-tailed field at the noisiest review; 0.14 to 0.18 in the heavy-tailed field with sharp to moderately noisy review.

a few researchers or spread across many. Our results so far rest on one production form, in which capability and resources are complements. One might hope that this choice settles the concentration question: the harder it is to substitute resources for capability, the more the money should go to the few researchers who can use it. To test this hope we embed our production form in a family of technologies indexed by an exponent γ , from Cobb-Douglas ($\gamma = 0$), under which a shortfall in either input can be offset by more of the other, through our form ($\gamma = -1$), to Leontief, under which output depends only on the scarcer input; Appendix C defines the family and reports which of our earlier results depend on the choice.²⁶ The technology does not settle the question (Figure 12). With a tight budget in a heavy-tailed field, harder complementarity does concentrate the informed funder’s grants. With an ample budget, harder complementarity spreads the informed funder’s grants, a pattern consistent with a funder that raises every researcher toward their bottleneck once the budget allows it. Where capability is evenly spread and the budget ample, the relation is not even monotone: concentration peaks at an intermediate complementarity and declines toward both Cobb-Douglas and Leontief.²⁷ How widely funds should be spread is set jointly by the technology, the budget, and the field’s capability distribution; no one of the three decides it alone.

Next we take stock: what the model implies for funders, and what its assumptions leave out.

9. DISCUSSION

We have compared the expected impact on research output of targeted funding to what it would be under equal division of the same budget, and to what it would be in the absence of any contribution by the funder. Our results isolate one quantity as key to strategies for optimal grant funding: the value of targeting researchers with the greatest gap between their capabilities and their resources. Two features of an area of research determine this value. The first is the funder’s budget relative to researchers’ own resources. Where the budget is small, dollars are scarce and where they go matters; where the budget can relieve most researchers’ bottlenecks,

²⁶ The concentration runs are single-round, with two hundred seeds per cell (four hundred at the refinement points of the evenly spread curve), and measure the review-informed funder’s grants, not the complete-information allocation. Leontief allocations are computed with a finer allocation step, and ties at the kink are broken by fill order.

²⁷ Gini coefficients of the informed funder’s grants, from Cobb-Douglas to Leontief: 0.92 to 0.95 in a heavy-tailed field with a tight budget ($\alpha_K = 1.3$, $b = 0.1$); 0.35 to 0.21 in the default field with an ample budget ($b = 1$); 0.63 to 0.45 in a heavy-tailed field with an ample budget. In the evenly spread field with an ample budget ($\alpha_K = 3.5$, $b = 1$) the Gini coefficient is 0.16 at Cobb-Douglas, peaks at 0.196 near $\gamma = -6$, and falls to 0.185 at $\gamma = -12$; the peak’s rise over the $\gamma = -12$ level is 3.5 standard errors.




FIGURE 12. Complementarity concentrates funding only when money is tight. The change in the Gini coefficient of the review-informed funder’s grants, relative to its Cobb-Douglas level, as the production technology moves from Cobb-Douglas ($\gamma = 0$) toward perfect complements (Leontief, detached marks at the left edge). With a tight budget in a heavy-tailed field, harder complementarity concentrates grants; with an ample budget it spreads them. Where capability is evenly spread and the budget ample, the relation is not even monotone: concentration peaks at an intermediate complementarity ($\gamma \approx -6$) and declines toward both ends. Cobb-Douglas levels: 0.92 (heavy-tailed, tight budget), 0.35 (default field, ample), 0.16 (evenly spread, ample). Two hundred seeds per cell (four hundred at the refinement points of the evenly spread curve).

the optimal allocation and equal division produce nearly the same output (§4, Figure 2). The second is how unequally capability is distributed. Where a few researchers can produce far more than the rest, getting resources to those few is most of what a funder can add (§6).

We showed how much various sources of information contribute to the effectiveness of targeting this gap. Researcher track records are modestly predictive, but they underdetermine whether a researcher’s productivity is bottlenecked by resources or by capability, and additional rounds of records improve the funder’s allocation only slowly (§5, Figure 4). Peer review supplies what records cannot: a signal of capability itself. The value of that signal increases with the two features that set the value of targeting, and with the signal’s informativeness. In a field where capability is heavy-tailed, most of the signal’s value comes from separating the few exceptional researchers from the rest, and that separation requires little information, so even coarse review recovers most of the value (§6, Figure 5). The size of grants also determines what a funder can learn. On thin grants, a

researcher's output is approximately proportional to their resources and nearly independent of their capability, so thin grants reveal almost nothing about who is capable; grants comparable to capability are what make records informative (§5, Figure 3). Finally, a funder that spends later, after observing output, allocates with more information than one that spends at once. The gain from spending later is real, but it is smaller than the gain from the review signal at every horizon and compounding rate we sweep (§7, Figure 9). Choosing whom to fund outweighs choosing when.

We also measured what a funder loses by forgoing targeting. Uniform seed grants and lotteries both give up part of the value of targeting on purpose, so both lose most output where that value is largest: tight budgets and heavy-tailed capability (§8, Figures 10 and 11). A lottery over any pool of researchers produces less expected output than an equal division of the same money over the same pool, because a certain grant produces more expected output than a chance at a larger grant of the same average size (Proposition 2). Where capability is evenly spread, the model's answer is therefore thin grants to many researchers, not a lottery among a fundable few. Whether funding should be concentrated on a few researchers or spread across many depends on the same two features and on the production technology, and the budget reverses the technology's effect: harder complementarity concentrates grants when the budget is tight in a heavy-tailed field and spreads them when the budget is ample (Figure 12).

These results have a practical consequence for funders. No funding mechanism has a verdict of its own. Peer review, seed grants, and lotteries each receive opposite verdicts in different fields, and the two features of the field decide which verdict applies. A funder choosing among mechanisms should therefore begin by estimating the two features for its own field. The funder's budget relative to researchers' total resources is the funder's own arithmetic. The inequality of capability is harder to observe. The distribution of output is the natural proxy for the distribution of capability, but where grants are thin, output tracks resources rather than capability, so the output distribution understates capability inequality (§5). Two cases show what follows. A small foundation funding a heavy-tailed field faces the largest value of targeting per dollar. For this foundation, review directed at identifying the exceptional few is worth most, and coarse review recovers most of that value; seed grants lose the most output; and a lottery among the fundable loses about a sixth of what selection by review gains. A large agency funding an evenly spread field faces a small value of targeting. For this agency, review of any informativeness adds little; seed grants and lotteries lose little; and the best use of the budget is thin grants to many researchers.

Three further implications concern peer review. The first concerns where review effort should go. Since most of review's value comes from separating the exceptional few from the rest, and fine distinctions among the many applications of comparable merit add little output, a review system that spends most of its effort ranking the middle of the distribution spends that effort where the model locates the least output (§6). The second concerns how to read the evidence on review's efficacy. The finding that percentile scores barely predict productivity comes from funded NIH grants with scores in the top fifth (Fang et al. 2016), and the studies that find modest predictive power also use funded grants (Li and Agha 2015). Both kinds of study measure prediction after review has already made the separation the model values, inside the band where the model expects review to add least. A weak association within that band is consistent with review of considerable value over the whole applicant pool. For the same reason, the mapping in §6 from measured accuracy to signal noise, which treated the measured accuracy as review's accuracy over the whole pool, is conservative. The third concerns calibration. A funder that treats review as sharper than it is loses more output than a funder that treats review as noisier than it is (§6, Figure 6). A funder uncertain about review's accuracy should therefore err toward undertrust.

The model holds several things fixed, and its results should be read within those limits. There is one funder. Competition among funders, and the process by which past funding attracts future funding across funders, lie outside the model, though the gap rule itself directs grants away from researchers whose resources are already large. Researchers do not respond strategically. Proposal effort, risk choice, and the cost of review lie outside the model; the contest models cited in §2 cover that ground. Output is a scalar rate. Projects have no heterogeneity and exploration has no value, so results about funding unexplored regions of a field have no counterpart here. The review signal is unbiased noise. Bias and conservatism in review, which are among the strongest arguments for lotteries, are not modeled. There is no entry, exit, or career structure. Grants are consumed and inputs are complements; the timing results depend on complementarity and the information results do not (Appendix C). The funder is an idealization that knows the model and updates correctly, so the values of information reported here are what such a funder extracts, not what a real funder would extract. Finally, the magnitudes are orderings and comparative statics over the parameters we sweep, and none is calibrated to a real funder. What a reader should carry to a real field is the pair of features that organizes the results, not the percentages.

The model makes three predictions that data on funded research can test. First, capability proxies, whether prior record or review score, should predict output

better the larger a grant is relative to the recipient's other resources, because on thin grants output is approximately proportional to resources and nearly independent of capability (§5, Figure 3). A test compares the predictive power of review scores across grant sizes within one applicant pool. Second, review's predictive validity, measured over the whole applicant pool rather than among funded grants, should be higher in fields whose output distributions have heavier tails (§6). A test compares fields, or one funder's programs, using the full range of scores where the outcomes of funded and unfunded applicants can both be observed. Third, where funders run lotteries beside review, the output lost to the lottery should be larger in heavier-tailed fields (§8, Figure 11). The lottery programs now running at several funders supply the comparison. Each prediction is comparative, and each can fail.

Appendices

A. DERIVATION OF THE GAP RULE

The within-round allocation problem. A population of n researchers. Researcher i has epistemic capability K_i and baseline resources R_{i0} , both fixed within the round. A grant $g_i \geq 0$ adds to researcher i 's resources for the round, so that their expected output is $A\Lambda(K_i, R_{i0} + g_i)$ with $\Lambda(K, R) = 2KR/(K + R)$, the production function of §3, and $A > 0$ the common productivity scale. Write

$$\varphi_i(g) = \frac{2AK_i(R_{i0} + g)}{K_i + R_{i0} + g}$$

for researcher i 's expected output as a function of their grant alone.

Throughout: $A > 0$; $K_i > 0$ and $R_{i0} \geq 0$ for every i ; and the round's tranche $B > 0$. (Under the model's continuous distributions for K and R , the conditions on K_i and R_{i0} hold with probability one.) Nothing else about the model is used: the proof does not depend on the distributions of capability and resources, on the noise in realized output, or on the dynamics across rounds. In particular, the result holds for every realized population and every tranche, not merely on average.

The funder observes $(K_i, R_{i0})_{i=1}^n$ and divides the tranche so as to maximize expected output:

$$(P) \quad \max_{g_1, \dots, g_n \geq 0} \sum_{i=1}^n \varphi_i(g_i) \quad \text{subject to} \quad \sum_{i=1}^n g_i \leq B.$$

Three lemmas.

Lemma 1 (the return to a grant). *For each i , on $g \in [0, \infty)$: (i) φ_i is strictly increasing, with $\varphi'_i(g) = 2AK_i^2/(K_i + R_{i0} + g)^2 > 0$; (ii) φ_i is strictly concave: $\varphi''_i(g) = -4AK_i^2/(K_i + R_{i0} + g)^3 < 0$; (iii) $\varphi'_i(g) < 2A$ whenever $R_{i0} + g > 0$, and $\varphi_i(g) \uparrow 2AK_i$ as $g \rightarrow \infty$.*

Proof. Write $S = K_i + R_{i0} + g$. Then $\varphi_i(g) = 2AK_i(S - K_i)/S = 2AK_i - 2AK_i^2/S$, and (i) and (ii) follow by differentiating in g (note $dS/dg = 1$). For (iii): $\varphi'_i(g) = 2AK_i^2/S^2 < 2A$ exactly when $S > K_i$, that is, when $R_{i0} + g > 0$; and $\varphi_i(g) = 2AK_i - 2AK_i^2/S \rightarrow 2AK_i$ as $S \rightarrow \infty$. $\square$

Lemma 1 states the economics of the problem. The marginal value of a dollar to researcher i , $\varphi'_i(g)$, increases with their capability and decreases with the resources they already hold, grants included. Each dollar granted lowers the value of the next (strict concavity). And expected output increases with every dollar but never exceeds $2AK_i$: it is bounded above by a quantity set by capability, which it approaches as $g \rightarrow \infty$ but does not attain.

Lemma 2 (existence, uniqueness, exhaustion). *Problem (P) has exactly one solution $g^* = (g_1^*, \dots, g_n^*)$, and it exhausts the tranche: $\sum_i g_i^* = B$.*

Proof. The feasible set $F = \{g \geq 0 : \sum_i g_i \leq B\}$ is compact and convex, and the objective $f(g) = \sum_i \varphi_i(g_i)$ is continuous, so a maximizer exists. f is strictly concave on F : for distinct feasible $g \neq h$ and $t \in (0, 1)$, concavity of each φ_i gives $f(tg + (1-t)h) \geq tf(g) + (1-t)f(h)$, and the inequality is strict because $g_j \neq h_j$ for some j and φ_j is strictly concave. If two distinct maximizers existed, their midpoint would be feasible (convexity of F) and strictly better, a contradiction; so the maximizer is unique. Finally, if $\sum_i g_i^* < B$, then increasing any one g_i^* slightly stays feasible and raises f (lemma 1(i)), contradicting optimality; so $\sum_i g_i^* = B$. $\square$

Lemma 3 (equal marginal value). *A feasible allocation g with $\sum_i g_i = B$ solves (P) if and only if there is a $\nu > 0$ such that*

$$\varphi'_i(g_i) = \nu \text{ for every } i \text{ with } g_i > 0, \quad \varphi'_i(0) \leq \nu \text{ for every } i \text{ with } g_i = 0.$$

Proof. Necessity. Let g^* solve (P). Since $B > 0$ is exhausted (lemma 2), some researcher is funded: $g_i^* > 0$ for some i . Fix any such i and any $j \neq i$, and for $0 < \delta \leq g_i^*$ consider the transfer that replaces g_i^* by $g_i^* - \delta$ and g_j^* by $g_j^* + \delta$, leaving the rest unchanged. The result is feasible, so optimality gives

$$\varphi_j(g_j^* + \delta) - \varphi_j(g_j^*) \leq \varphi_i(g_i^*) - \varphi_i(g_i^* - \delta).$$

Dividing by δ and letting $\delta \downarrow 0$ yields $\varphi'_j(g_j^*) \leq \varphi'_i(g_i^*)$. If j is also funded, the same argument with i and j exchanged gives the reverse inequality, so all funded researchers share a common marginal value; call it $\nu = \varphi'_i(g_i^*) > 0$ (positivity by lemma 1(i)). For unfunded j , the displayed inequality reads $\varphi'_j(0) \leq \nu$.

Sufficiency. Let g satisfy the condition with multiplier ν , and let h be any feasible allocation. Concavity of φ_i gives, for every i , $\varphi_i(h_i) \leq \varphi_i(g_i) + \varphi'_i(g_i)(h_i - g_i)$. For funded i , $\varphi'_i(g_i) = \nu$, so $\varphi'_i(g_i)(h_i - g_i) = \nu(h_i - g_i)$. For unfunded i , $g_i = 0$ and $h_i - g_i = h_i \geq 0$, so $\varphi'_i(0)(h_i - g_i) \leq \nu(h_i - g_i)$. Summing over i :

$$f(h) \leq f(g) + \nu \sum_i (h_i - g_i) = f(g) + \nu \left(\sum_i h_i - B \right) \leq f(g),$$

since h is feasible and $\nu > 0$. So g solves (P). $\square$

Lemma 3 is the formal content of the sentence in §4: an allocation is optimal exactly when no dollar can be moved to a researcher who can produce more value with it.

The proposition.

Proposition 1 (the gap rule). *The unique solution of (P) is*

$$g_i^* = \max(c K_i - R_{i0}, 0), \quad i = 1, \dots, n,$$

where $c > 0$ is the unique constant satisfying the budget identity $\sum_{i=1}^n \max(c K_i - R_{i0}, 0) = B$. Equivalently, $c = \sqrt{2A/\nu} - 1$, where $\nu \in (0, 2A)$ is the common marginal value of lemma 3.

Proof. Let g^* be the unique solution (lemma 2) and ν the multiplier that lemma 3 associates with it. The proof proceeds in four steps.

Step 1: $\nu \in (0, 2A)$, so $c = \sqrt{2A/\nu} - 1$ is well defined and positive. Positivity of ν is part of lemma 3. For the upper bound: some researcher i is funded, and for them $R_{i0} + g_i^* \geq g_i^* > 0$, so lemma 1(iii) gives $\nu = \varphi'_i(g_i^*) < 2A$. Then $2A/\nu > 1$, so $c = \sqrt{2A/\nu} - 1 > 0$.

Step 2: funded researchers receive exactly their gap. Let $g_i^* > 0$. By lemma 3, $2A K_i^2 / (K_i + R_{i0} + g_i^*)^2 = \nu$. Solving, and taking the positive square root (both sides of the next display are positive),

$$K_i + R_{i0} + g_i^* = K_i \sqrt{2A/\nu} = (1 + c) K_i,$$

so $g_i^* = c K_i - R_{i0}$. Since $g_i^* > 0$, this equals $\max(c K_i - R_{i0}, 0)$.

Step 3: unfunded researchers have no gap. Let $g_i^* = 0$. By lemma 3, $2AK_i^2/(K_i + R_{i0})^2 = \varphi_i'(0) \leq \nu$, so $K_i + R_{i0} \geq K_i \sqrt{2A/\nu} = (1+c)K_i$, that is, $R_{i0} \geq cK_i$. Then $\max(cK_i - R_{i0}, 0) = 0 = g_i^*$.

Steps 2 and 3 establish $g_i^* = \max(cK_i - R_{i0}, 0)$ for every i , and the budget identity is exhaustion (lemma 2). It remains to show c is the only constant satisfying that identity.

Step 4: the budget identity pins c uniquely. Define $G(c) = \sum_{i=1}^n \max(cK_i - R_{i0}, 0)$ for $c \geq 0$. G is continuous and nondecreasing, with $G(0) = 0$ and $G(c) \rightarrow \infty$ as $c \rightarrow \infty$. Moreover G is strictly increasing wherever it is positive: if $G(c) > 0$, then $cK_j > R_{j0}$ for some j , and for any $c' > c$, $G(c') \geq G(c) + K_j(c' - c) > G(c)$. Now suppose $c_1 < c_2$ both satisfy $G(c) = B$. Since $B > 0$, $G(c_1) > 0$, so $G(c_2) > G(c_1) = B$, a contradiction. $\square$

Consequences.

Corollary 1 (funding frontier and budget comparative statics). *(i) Researcher i is funded, $g_i^* > 0$, exactly when $K_i > R_{i0}/c$. (ii) As a function of the tranche B , the constant $c(B)$ is continuous and strictly increasing; hence the funded set is nondecreasing in B , every funded researcher's grant $c(B)K_i - R_{i0}$ is strictly increasing in B , and the common marginal value $\nu(B) = 2A/(1+c(B))^2$ is strictly decreasing in B .*

Proof. (i) is immediate from the formula: $\max(cK_i - R_{i0}, 0) > 0$ iff $cK_i > R_{i0}$. For (ii): G is continuous, strictly increasing where positive, and unbounded, so $c(B) = G^{-1}(B)$ is well defined, continuous, and strictly increasing on $B > 0$ (Step 4). The remaining claims follow by monotonicity of $c \mapsto cK_i - R_{i0}$ and $c \mapsto 2A/(1+c)^2$. $\square$

Corollary 2 (the value of targeting vanishes in the budget). *Let $f^*(B)$ denote the optimal value of (P) and $f^u(B) = \sum_i \varphi_i(B/n)$ the value of uniform funding. Then $0 \leq f^*(B) - f^u(B)$ for every B , and $f^*(B) - f^u(B) \rightarrow 0$ as $B \rightarrow \infty$.*

Proof. The first inequality holds because uniform funding is feasible. For the second: by lemma 1(iii), $\varphi_i(g) < 2AK_i$ for every g , so $f^*(B) \leq 2A \sum_i K_i$; and $f^u(B) = \sum_i \varphi_i(B/n) \rightarrow 2A \sum_i K_i$ as $B \rightarrow \infty$. Hence $0 \leq f^*(B) - f^u(B) \leq 2A \sum_i K_i - f^u(B) \rightarrow 0$. $\square$

As the budget increases, any allocation that spreads funds widely brings every researcher's output near its upper bound $2AK_i$, so allocations come to differ little in total output. Since uniform funding's own gain over no funding does not vanish

(it approaches $2A \sum_i K_i - \sum_i \varphi_i(0) > 0$), the same conclusion holds for the value of targeting measured relative to that gain.

Example 1 (funding the track record can produce less than uniform funding). *Two researchers, $A = 1$, tranche $B = 2$. Researcher 1: $K_1 = 8$, $R_{10} = 8$ (productive and well resourced; expected output $\varphi_1(0) = 8$). Researcher 2: $K_2 = 4$, $R_{20} = 1$ (capable but resource-starved; $\varphi_2(0) = 8/5$). Allocating in proportion to expected output gives $g = (5/3, 1/3)$ and total output $570/53 \approx 10.75$. Uniform funding, $g = (1, 1)$, gives $568/51 \approx 11.14$. The gap rule ($c = 3/4$) gives $g^* = (0, 2)$ and $80/7 \approx 11.43$. Track-record funding sends most of the tranche to the researcher with the small gap and produces less output than uniform funding.*

Example 2 (funding the under-resourced can produce less than uniform funding). *Two researchers, $A = 1$, tranche $B = 2$. Researcher 1: $K_1 = 10$, $R_{10} = 5$. Researcher 2: $K_2 = 1/2$, $R_{20} = 0$ (the scarcest resources in the population). Directing the tranche to the most resource-starved researcher gives total output $112/15 \approx 7.47$. Uniform funding gives $49/6 \approx 8.17$. The gap rule ($c = 2/3$) gives $g^* = (5/3, 1/3)$ and $42/5 = 8.4$. Scarcity alone misdirects the tranche toward low capability; note that the optimum still grants the scarce researcher a little, because their marginal value at zero is high, but that marginal value is determined by capability, not by scarcity.*

Remark 1 (water-filling). *At the optimum, every funded researcher's total resources are filled to the capability-scaled level $R_{i0} + g_i^* = c K_i$, equivalently $K_i + R_{i0} + g_i^* = (1 + c)K_i$. Funding to a common marginal value by filling gaps from a level is the water-filling solution of information theory (Cover and Thomas 2006); here the fill level scales with capability rather than being common to all.*

Remark 2 (the capability-independent special case). *If all researchers share one capability, $K_i = K$, the fill level cK is common, and the rule reduces to filling each researcher's gap below a single target: the poverty-gap transfer of Foster et al. (1984).*

Remark 3 (scope across rounds). *Proposition 1 governs the division of a given tranche within a round, taking that round's (K_i, R_{i0}) as fixed. Across rounds, capability compounds and the constant c and the funded set are recomputed from the new population state; nothing in the proof depends on how the population state arose. When the funder also chooses how to split the budget across rounds, the proposition governs the within-round division of whatever tranche that choice assigns.*

B. LOTTERIES AND EQUAL DIVISION

Fix a round, a budget $B > 0$, and a pool $P \subseteq \{1, \dots, n\}$ of m researchers. The *equal division* over P grants each researcher in P the amount B/m . A *lottery* over P with $k \leq m$ winners selects k researchers from P uniformly at random and grants each winner B/k . Expected output is evaluated as in Appendix A: given an allocation g , it is $\sum_i \varphi_i(g_i)$, and for a random allocation, the expectation is taken over the draw as well.

Proposition 2 (lotteries are outperformed by equal division). *For every pool P , budget B , and winner count $k \leq m$, the lottery's expected output does not exceed the equal division's, with equality only if $k = m$. In particular, a lottery over the whole population, at any winner count, produces at most uniform funding's expected output.*

Proof. Researchers outside P receive nothing under either scheme. For $i \in P$, the lottery's grant to i is a random variable $\tilde{g}_i$ equal to B/k with probability k/m and 0 otherwise, so $\mathbb{E}[\tilde{g}_i] = B/m$. By linearity, the lottery's expected output is $\sum_{i \notin P} \varphi_i(0) + \sum_{i \in P} \mathbb{E}[\varphi_i(\tilde{g}_i)]$. Since φ_i is strictly concave (Lemma 1(ii)), Jensen's inequality gives $\mathbb{E}[\varphi_i(\tilde{g}_i)] \leq \varphi_i(\mathbb{E}[\tilde{g}_i]) = \varphi_i(B/m)$ for every $i \in P$, with equality only if $\tilde{g}_i$ is degenerate, that is, $k = m$. Summing over i yields the claim; taking $P = \{1, \dots, n\}$ yields the comparison with uniform funding. $\square$

Remark 4 (what the proposition does and does not cover). *The proposition compares expected output within a fixed pool. It measures the effect of the randomization alone: whatever pool review selects, replacing the lottery within that pool by the equal division of the same money weakly increases expected output. The proposition says nothing about the choice of pool, about the review ranking within the pool that the equal division also ignores, or about anything a lottery saves outside the model (review costs, proposal effort, bias). § 8 measures the first two and cites the literature on the third.*

C. ROBUSTNESS TO THE PRODUCTION TECHNOLOGY

Throughout the paper, a researcher's expected output has taken the form of §3, $2AKR/(K + R)$, in which capability and resources are complements and output is limited by the scarcer of the two. This appendix asks which of our results depend on that choice. We embed our form in a family of production technologies indexed by an exponent γ : expected output is $AM_\gamma(K, R)$, where $M_\gamma(K, R) = ((K^\gamma + R^\gamma)/2)^{1/\gamma}$ is the power mean of capability and resources. At $\gamma = -1$ the power mean is the harmonic mean $2KR/(K + R)$, and the family reproduces our form. In the limit $\gamma \rightarrow 0$, the Cobb-Douglas case, the power mean is the geometric

mean $\sqrt{KR}$, and a shortfall in either input can be offset by more of the other; as γ decreases the inputs grow harder to substitute for one another; and in the limit $\gamma \rightarrow -\infty$, the Leontief case, the power mean is $\min(K, R)$, and output depends only on the scarcer input. We re-run the model’s central comparisons across this family.²⁸

The value of the review signal survives the whole family (Figure 13). At every technology tested, from Cobb-Douglas to Leontief, the signal’s value increases as capability grows more heavy-tailed and as review grows more informative. The ordering of review’s value by field (§6) is therefore not an artifact of our production form. What changes with the technology is the scale. The more complementary the inputs, the more review is worth: under heavy-tailed capability with sharp review, the signal’s value increases from about seven percent of no-funding output at Cobb-Douglas to about thirty percent at $\gamma = -3$ and a third at Leontief.²⁹ Complementarity raises what targeting is worth, and with it the value of the information that targeting requires.

The timing results of §7 do not survive the substitutable end of the family (Table 2). Under Cobb-Douglas production the value of scheduling is indistinguishable from zero at every compounding rate, and the optimal schedule stays within about a hundredth of even. At $\gamma = -3$, a technology more complementary than our form, both the value of scheduling and the lean toward late spending increase with the compounding rate and exceed their Cobb-Douglas values at every rate. The timing results hold where capability and resources are complements and vanish where they are substitutable.

How widely the informed funder spreads its grants across this family is reported in §8 (Figure 12).

REFERENCES

- Kaare Aagaard, Alexander Kladakis, and Mathias W. Nielsen. Concentration or dispersal of research funding? *Quantitative Science Studies*, 1(1):117–149, 2020.
- Shahar Avin. Mavericks and lotteries. *Studies in History and Philosophy of Science Part A*, 76:13–23, 2019a.
- Shahar Avin. Centralized funding and epistemic exploration. *The British Journal for the Philosophy of Science*, 70(3):629–656, 2019b.

²⁸ Two rounds, one hundred seeds per cell (fifty for Leontief). Leontief allocations are computed with a finer allocation step, and ties at the kink are broken by fill order; the Leontief points carry that caveat. These runs use an earlier version of the allocation routine than the runs reported in the body, so their magnitudes are compared with one another and not with the body’s figures.

²⁹ 6.8, 29.3, and 33.7 percent of no-funding output; the same ordering holds at every tested tail and noise level at which the value is distinguishable from zero.




FIGURE 13. The structure of review’s value survives the production technology. The value of the review signal as a percent of no-funding output, across the CES family: at every technology tested, the value increases as capability grows more heavy-tailed (left, decreasing α_K) and as review grows more informative (right, decreasing τ_K), and its scale increases with complementarity. Filled circles: Cobb-Douglas; open circles: $\gamma = -3$; triangles: Leontief. Two rounds; one hundred seeds per point (fifty for Leontief).

TABLE 2. The value of scheduling and the optimal schedule by production technology. Five rounds. The value of scheduling is the forward-looking funder’s gain over the myopic funder, both with review, as a percent of no-funding output; the schedule’s center of mass is 0.5 when spending is even and larger when spending leans late. One hundred seeds per cell; standard errors of the gain are at most 0.002 (Cobb-Douglas) and 0.09 ($\gamma = -3$) percentage points.

compounding rate ϵ	0.05	0.15	0.30	0.55	0.85
<i>Value of scheduling (percent of no-funding output)</i>					
Cobb-Douglas ($\gamma = 0$)	-0.001	-0.002	0.001	0.001	0.004
$\gamma = -3$	0.01	0.13	0.45	1.07	1.54
<i>Center of mass of the optimal schedule</i>					
Cobb-Douglas ($\gamma = 0$)	0.499	0.500	0.501	0.505	0.512
$\gamma = -3$	0.511	0.540	0.577	0.623	0.660

Thijs Bol, Mathijs de Vaan, and Arnout van de Rijt. The Matthew effect in science funding. *Proceedings of the National Academy of Sciences*, 115(19):4887–4890, 2018.

Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. Wiley-Interscience, Hoboken, NJ, 2nd edition, 2006.

- Ferric C. Fang and Arturo Casadevall. Research funding: the case for a modified lottery. *mBio*, 7(2):e00422–16, 2016. fields consistent across secondary sources; confirm against DOI at submission.
- Ferric C. Fang, Anthony Bowen, and Arturo Casadevall. NIH peer review percentile scores are poorly predictive of grant productivity. *eLife*, 5:e13323, 2016.
- James Foster, Joel Greer, and Erik Thorbecke. A class of decomposable poverty measures. *Econometrica*, 52(3):761–766, 1984.
- Stephen A. Gallo, Afton S. Carpenter, David Irwin, Caitlin D. McPartland, Joseph Travis, Sofie Reynders, Lisa A. Thompson, and Scott R. Glisson. The validation of peer review through research impact measures and the implications for funding strategies. *PLoS ONE*, 9(9):e106474, 2014.
- Kevin Gross and Carl T. Bergstrom. Contest models highlight inherent inefficiencies of scientific funding competitions. *PLoS Biology*, 17(1):e3000065, 2019.
- Kevin Gross and Carl T. Bergstrom. How competition propels scientific risk-taking. arXiv:2509.06718, 2025.
- Remco Heesen. Cumulative advantage and the incentive to commit fraud in science. *The British Journal for the Philosophy of Science*, 75(3), 2024. doi 10.1086/716235; CHECK page range.
- Rachel Heyard, Manuela Ott, Georgia Salanti, and Matthias Egger. Rethinking the funding line at the Swiss National Science Foundation: Bayesian ranking and lottery. *Statistics and Public Policy*, 9(1):110–121, 2022.
- Yarden Katz and Ulrich Matter. Metrics of financial inequality and the concentration of research funding in the U.S. National Institutes of Health, 1985–2015. Working paper, arXiv:2007.09538. CHECK: published version may exist, 2020.
- Michael S. Lauer and Deepshikha Roychowdhury. Inequalities in the distribution of National Institutes of Health research project grant funding. *eLife*, 10:e71712, 2021.
- Danielle Li and Leila Agha. Big names or big ideas: Do peer-review panels select the best science proposals? *Science*, 348(6233):434–438, 2015.
- Mengyao Liu, Vernon Choy, Philip Clarke, Adrian Barnett, Tony Blakely, and Lucy Pomeroy. The acceptability of using a lottery to allocate research funding: a survey of applicants. *Research Integrity and Peer Review*, 5:3, 2020.
- Alfred J. Lotka. The frequency distribution of scientific productivity. *Journal of the Washington Academy of Sciences*, 16(12):317–323, 1926.
- Finn Luebber, Sören Krach, Frieder M. Paulus, Lena Rademacher, and Rima-Maria Rahal. Lottery before peer review is associated with increased female representation and reduced estimated economic cost in a German funding line. *Nature Communications*, 16:9824, 2025.

- Philippe Mongeon, Christine Brodeur, Catherine Beaudry, and Vincent Larivière. Concentration of research funding leads to decreasing marginal returns. *Research Evaluation*, 25(4):396–404, 2016.
- National Center for Science and Engineering Statistics. U.S. R&D totaled \$937 billion in 2023; estimate for 2024 indicates further increase to \$993 billion. Technical Report NSF 26-314, National Science Foundation, 2026. URL <https://ncses.nsf.gov/pubs/nsf26314>.
- National Science Board. The state of U.S. science and engineering 2026. Technical report, National Science Foundation, 2026. URL <https://ncses.nsf.gov/pubs/nsbsep20261>. CHECK report number.
- Paul Travis Nicholls. Bibliometric modeling processes and the empirical validity of Lotka’s law. *Journal of the American Society for Information Science*, 40(6):379–385, 1989. CHECK fields and exponent range.
- Hyunwoo Park, Jeongsik Lee, and Byung-Cheol Kim. Project selection in NIH: A natural experiment from ARRA. *Research Policy*, 44(6):1145–1159, 2015.
- Elizabeth L. Pier, Markus Brauer, Amarette Filut, Anna Kaatz, Joshua Raclaw, Mitchell J. Nathan, Cecilia E. Ford, and Molly Carnes. Low agreement among reviewers evaluating the same NIH grant applications. *Proceedings of the National Academy of Sciences*, 115(12):2952–2957, 2018.
- Lambros Roumbanis. Peer review or lottery? A critical analysis of two different forms of decision-making mechanisms for allocation of research grants. *Science, Technology, & Human Values*, 44(6):994–1019, 2019.
- Carlos Oscar S. Sorzano and B. Pueche-Granados. Research funding as a decision problem under heavy-tailed uncertainty. arXiv:2604.22793, 2026.